Advanced Data Engineering in Cloud

ASSIGNMENT-2

Data Ingestion and Processing

Hem Prakash Dev **Roll-G23AI1054**

1. **Implement the data ingestion mechanism using AWS Kinesis Data Streams or AWS Direct Connect to stream data from a source to Amazon S3.**

   a. **Create a Kinesis Data Stream**:

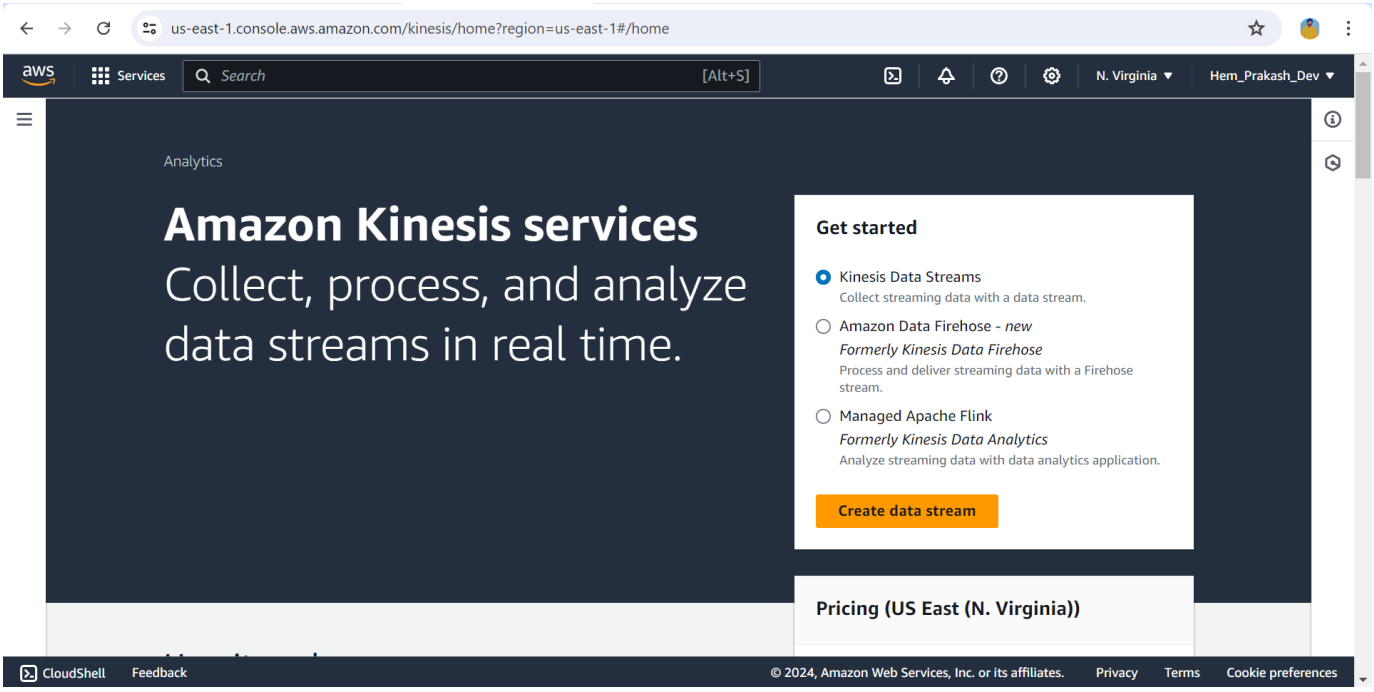

b.

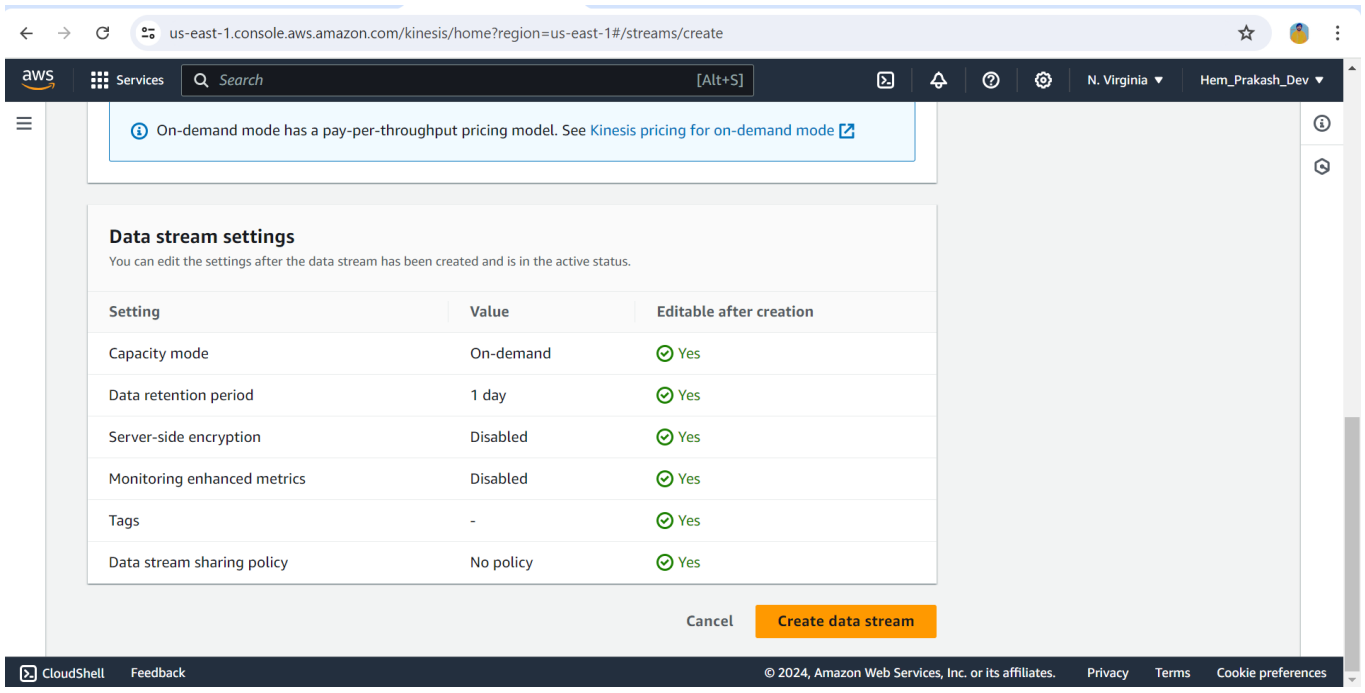

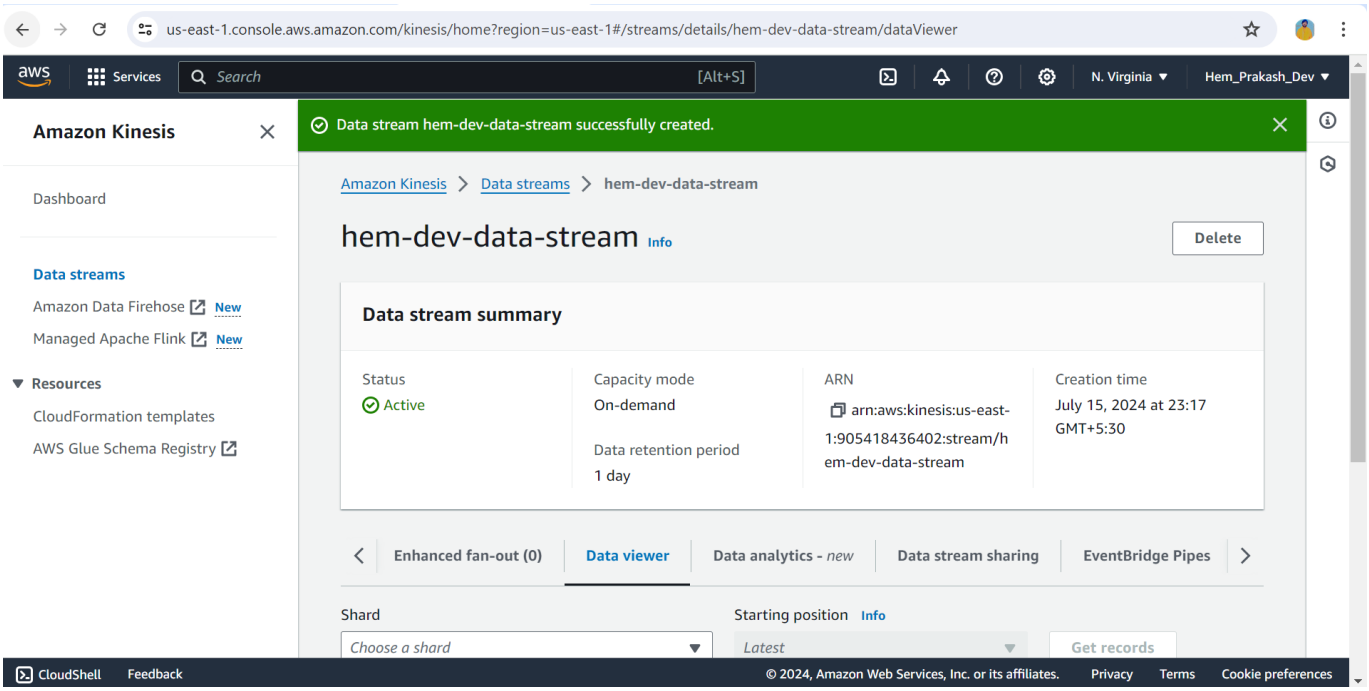

b. stream Data to S3 using Kinesis Firehose:

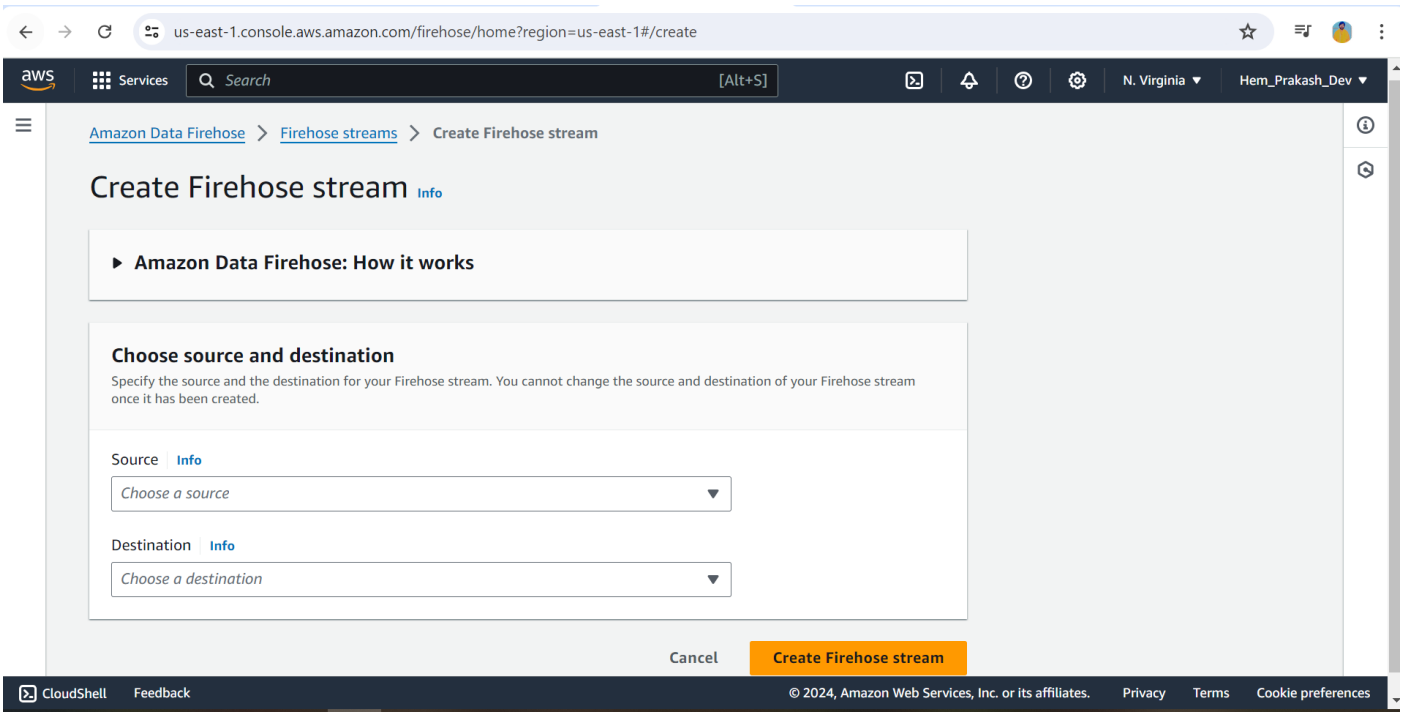

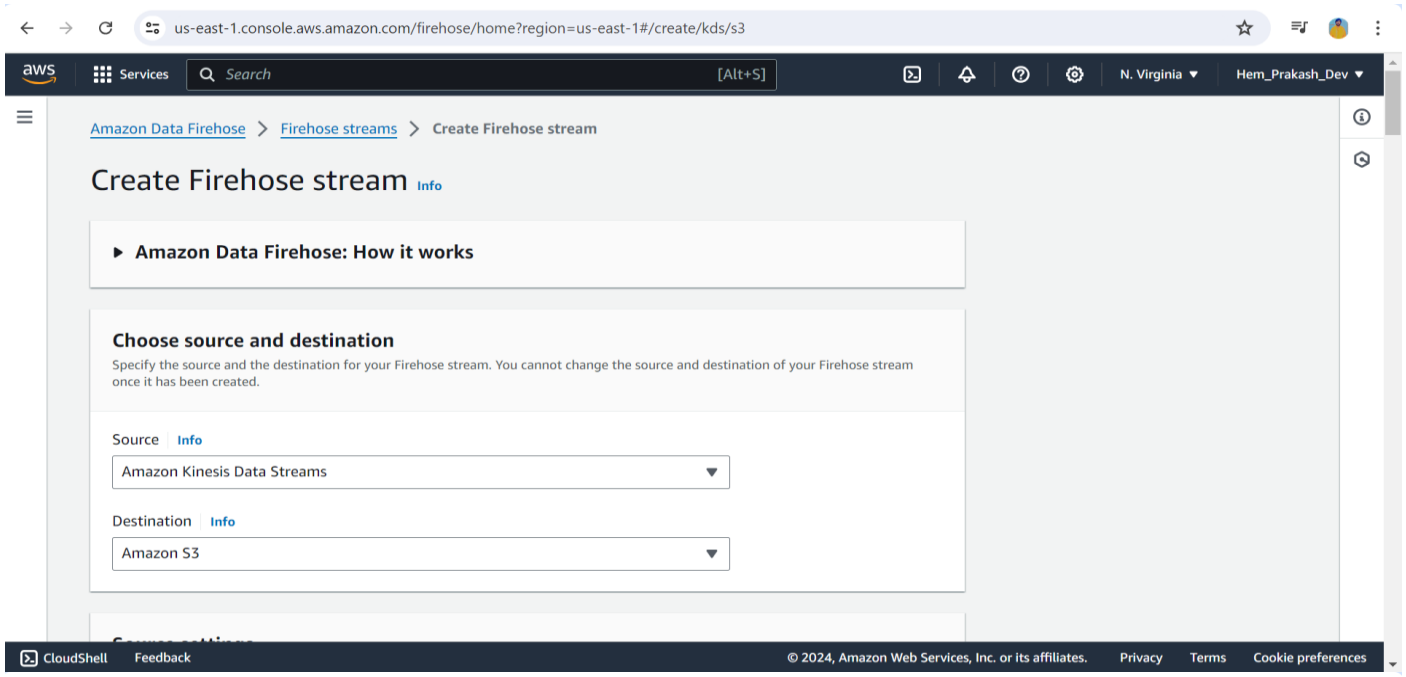
us-east-1.console.aws.amazon.com/firehose/home?region=us-east-1#/create/kds/s3
aws
Services
Search
[Alt+S]
N. Virginia
Hem_Prakash_Dev
Amazon Data Firehose > Firehose streams > Create Firehose stream
Create Firehose stream Info
Amazon Data Firehose: How it works
Choose source and destination
Specify the source and the destination for your Firehose stream. You cannot change the source and destination of your Firehose stream once it has been created.
Source Info
Amazon Kinesis Data Streams
Destination Info
Amazon S3
CloudShell
Feedback
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

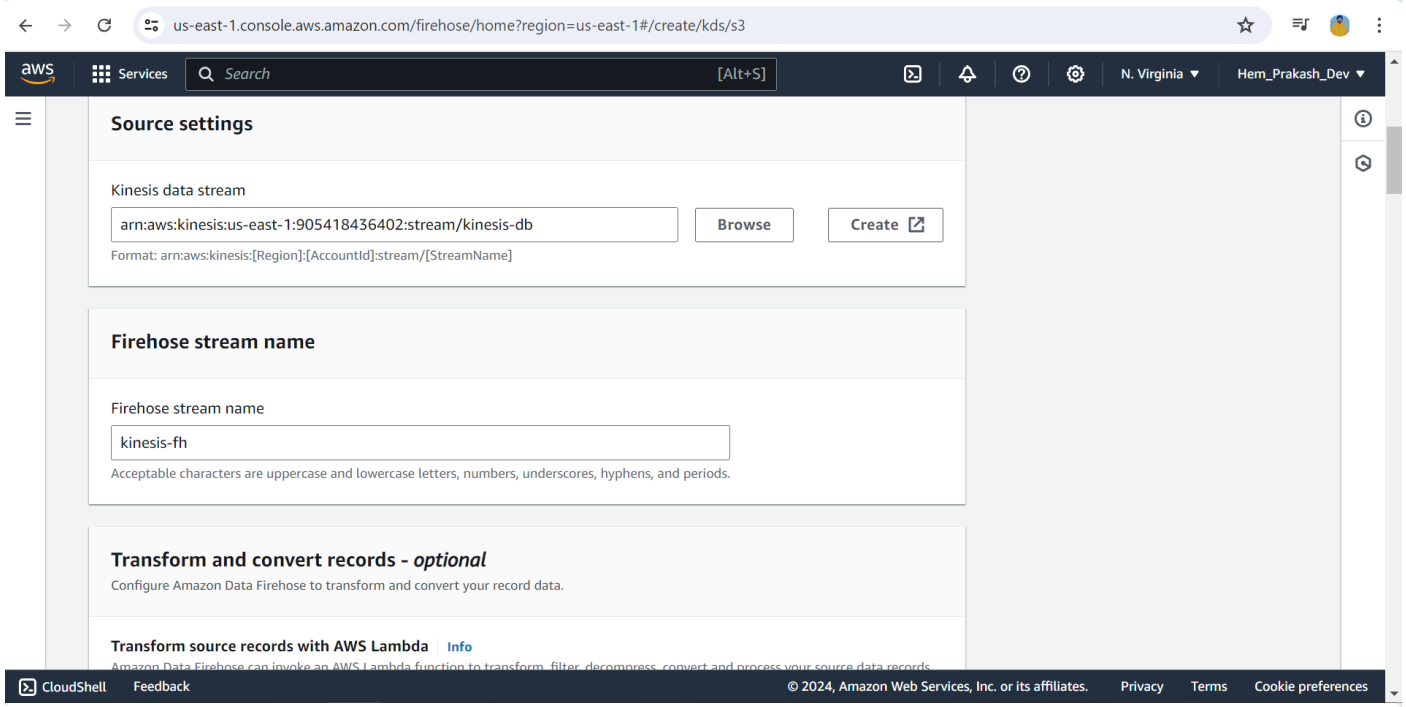
us-east-1.console.aws.amazon.com/firehose/home?region=us-east-1#/create/kds/s3
Source settings
Kinesis data stream
arn:aws:kinesis:us-east-1:905418436402:stream/kinesis-db
Browse
Create
Format: arn:aws:kinesis:[Region]:[AccountId]:stream/[StreamName]
Firehose stream name
Firehose stream name
kinesis-fh
Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens, and periods.
Transform and convert records - optional
Configure Amazon Data Firehose to transform and convert your record data.
Transform source records with AWS Lambda Info

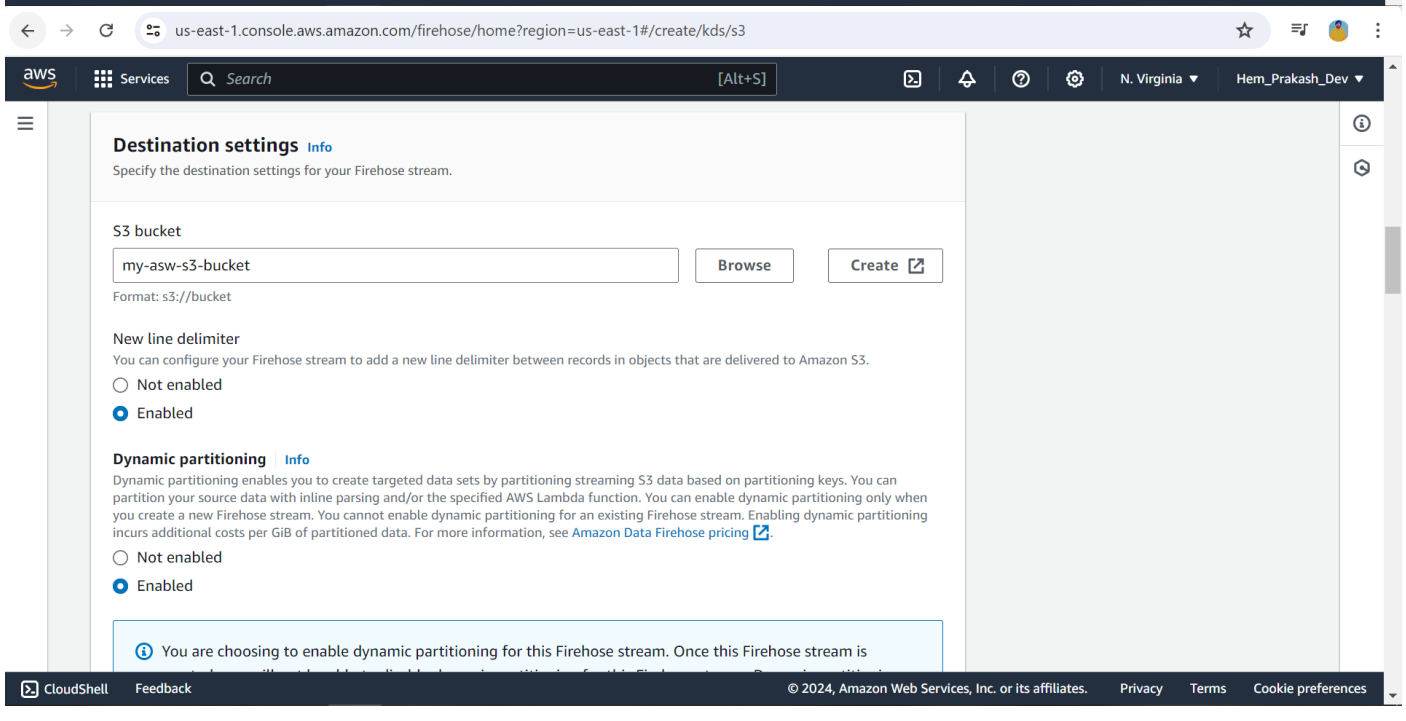
us-east-1.console.aws.amazon.com/firehose/home?region=us-east-1#/create/kds/s3
Destination settings Info
Specify the destination settings for your Firehose stream.
S3 bucket
my-asw-s3-bucket
Browse
Create
Format: s3://bucket
New line delimiter
You can configure your Firehose stream to add a new line delimiter between records in objects that are delivered to Amazon S3.
Not enabled
Enabled
Dynamic partitioning Info
Dynamic partitioning enables you to create targeted data sets by partitioning streaming S3 data based on partitioning keys. You can partition your source data with inline parsing and/or the specified AWS Lambda function. You can enable dynamic partitioning only when you create a new Firehose stream. You cannot enable dynamic partitioning for an existing Firehose stream. Enabling dynamic partitioning incurs additional costs per GiB of partitioned data. For more information, see Amazon Data Firehose pricing.
Not enabled
Enabled
You are choosing to enable dynamic partitioning for this Firehose stream. Once this Firehose stream is

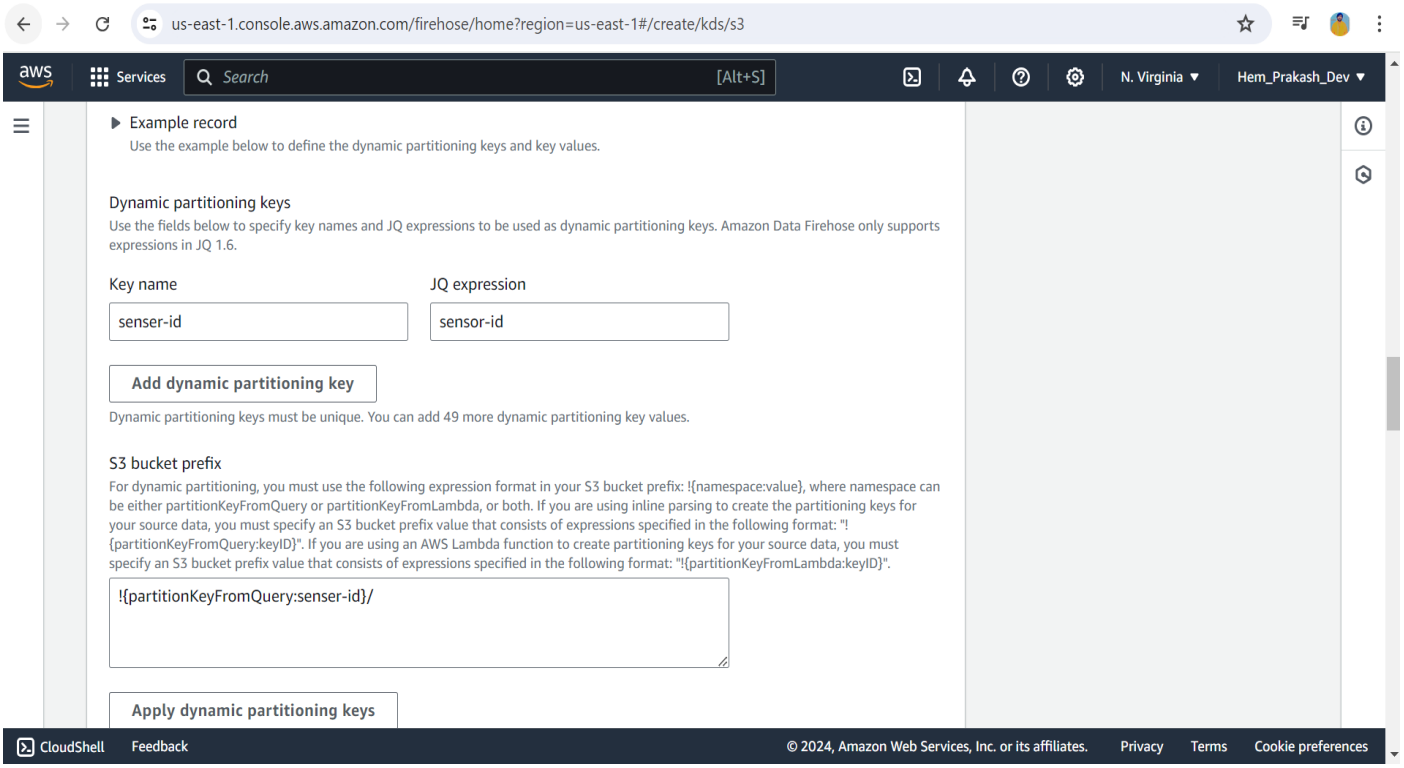

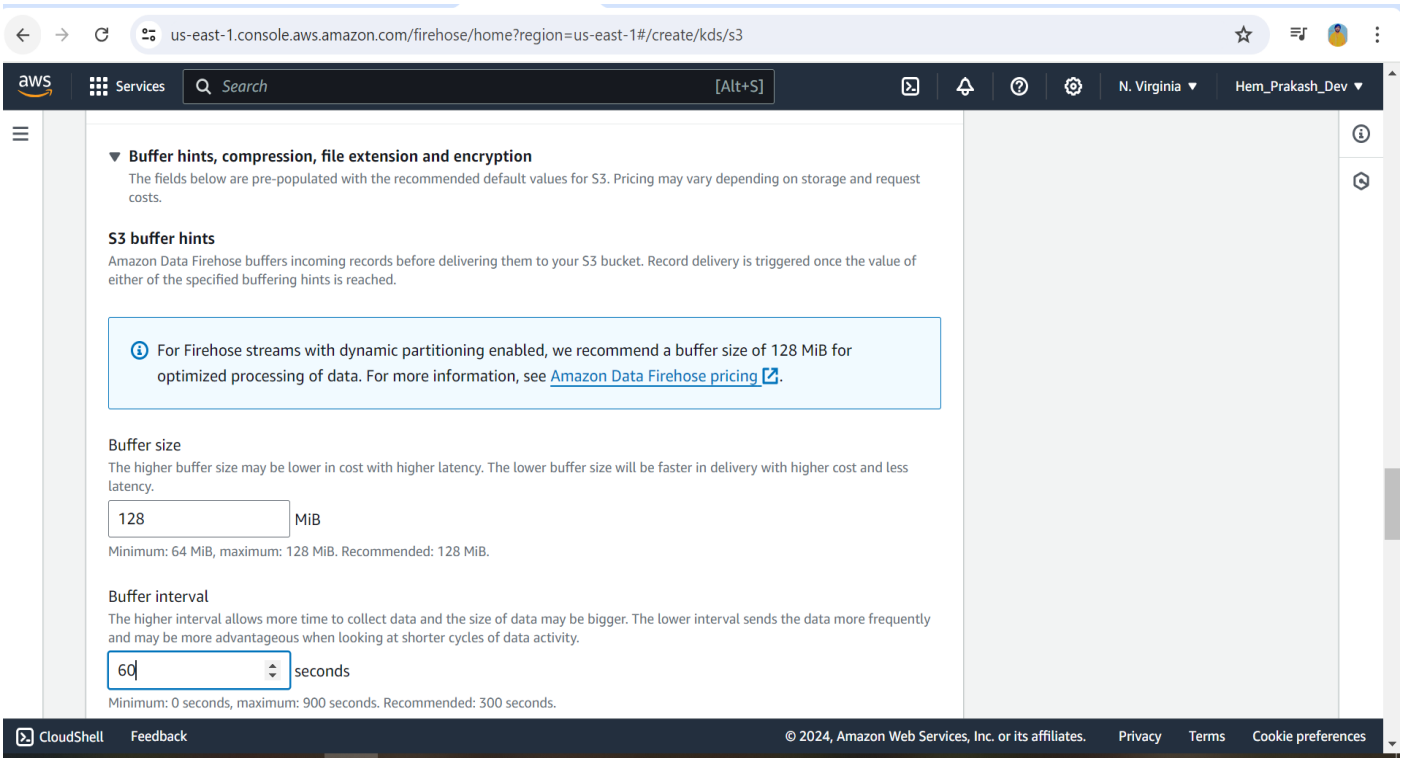

Kinesis firehose successfully created

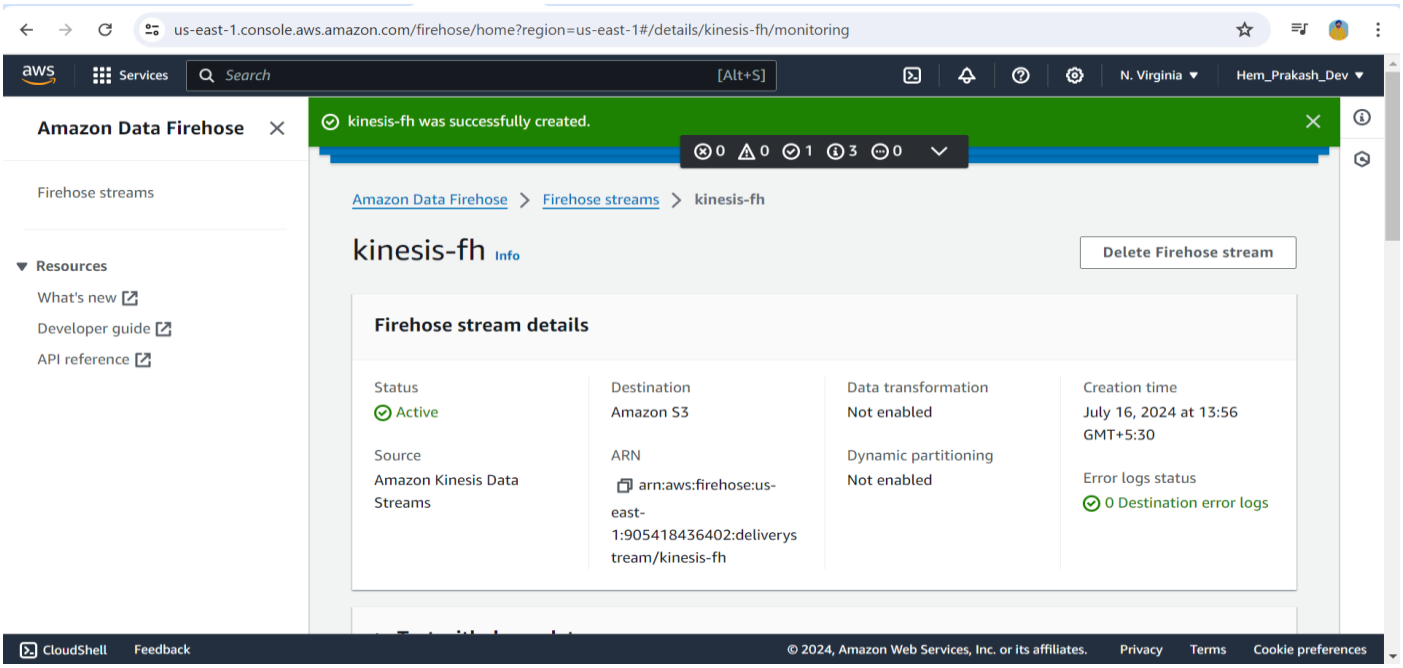

2. Develop and test the data processing pipeline using AWS Glue or Amazon EMR (Elastic MapReduce) with Apache Spark or Hadoop.

Data processing pipeline using Amazon EMR (Elastic MapReduce) with Apache Spark

EMR-Console

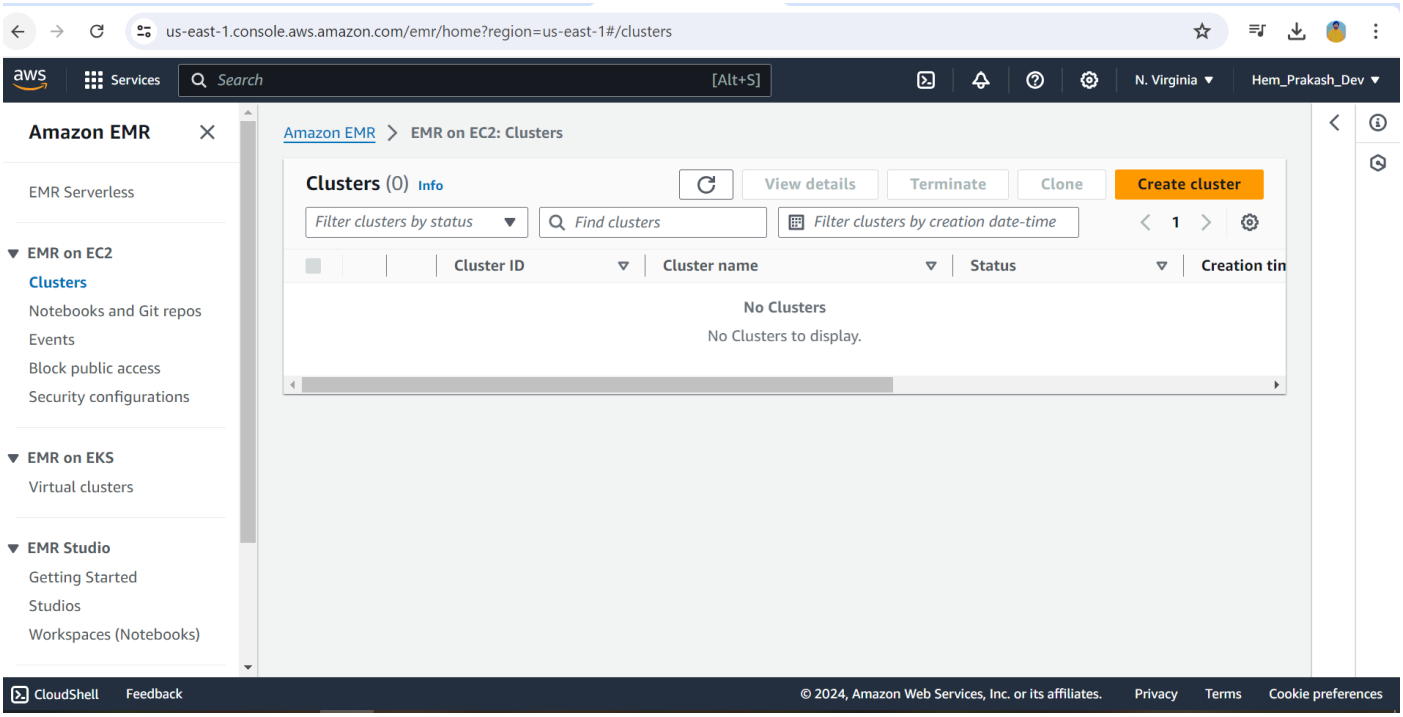

Cluster creation

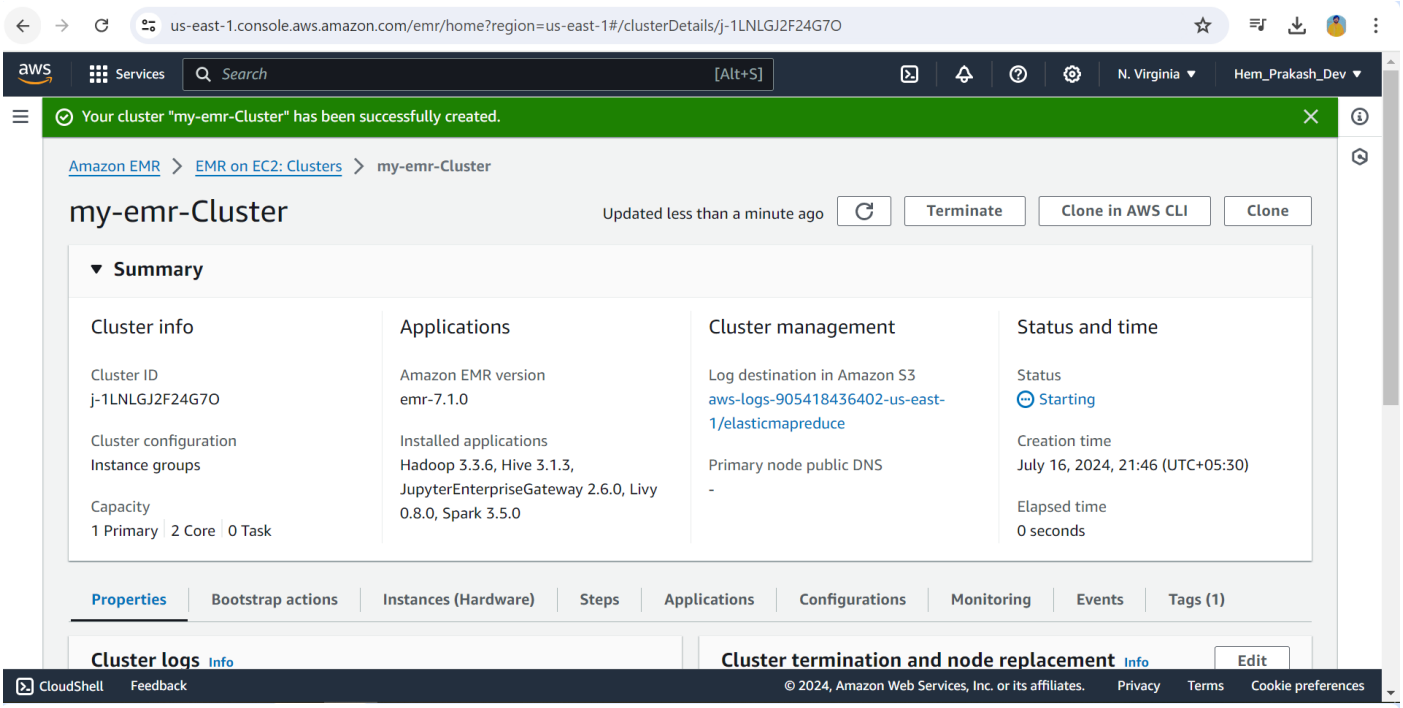
us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-1LNLGJ2F24G7O
Services
Search
[Alt+S]
N. Virginia
Hem_Prakash_Dev
Your cluster "my-emr-Cluster" has been successfully created.
Amazon EMR
EMR on EC2: Clusters
my-emr-Cluster
my-emr-Cluster
Updated less than a minute ago
Terminate
Clone in AWS CLI
Clone
Summary
Cluster info
Cluster ID
j-1LNLGJ2F24G7O
Cluster configuration
Instance groups
Capacity
1 Primary | 2 Core | 0 Task
Applications
Amazon EMR version
emr-7.1.0
Installed applications
Hadoop 3.3.6, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Livy 0.8.0, Spark 3.5.0
Cluster management
Log destination in Amazon S3
aws-logs-905418436402-us-east-1/elasticmapreduce
Primary node public DNS
-
Status and time
Status
Starting
Creation time
July 16, 2024, 21:46 (UTC+05:30)
Elapsed time
0 seconds
Properties
Bootstrap actions
Instances (Hardware)
Steps
Applications
Configurations
Monitoring
Events
Tags (1)
Cluster logs Info
Cluster termination and node replacement Info
Edit
CloudShell
Feedback
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

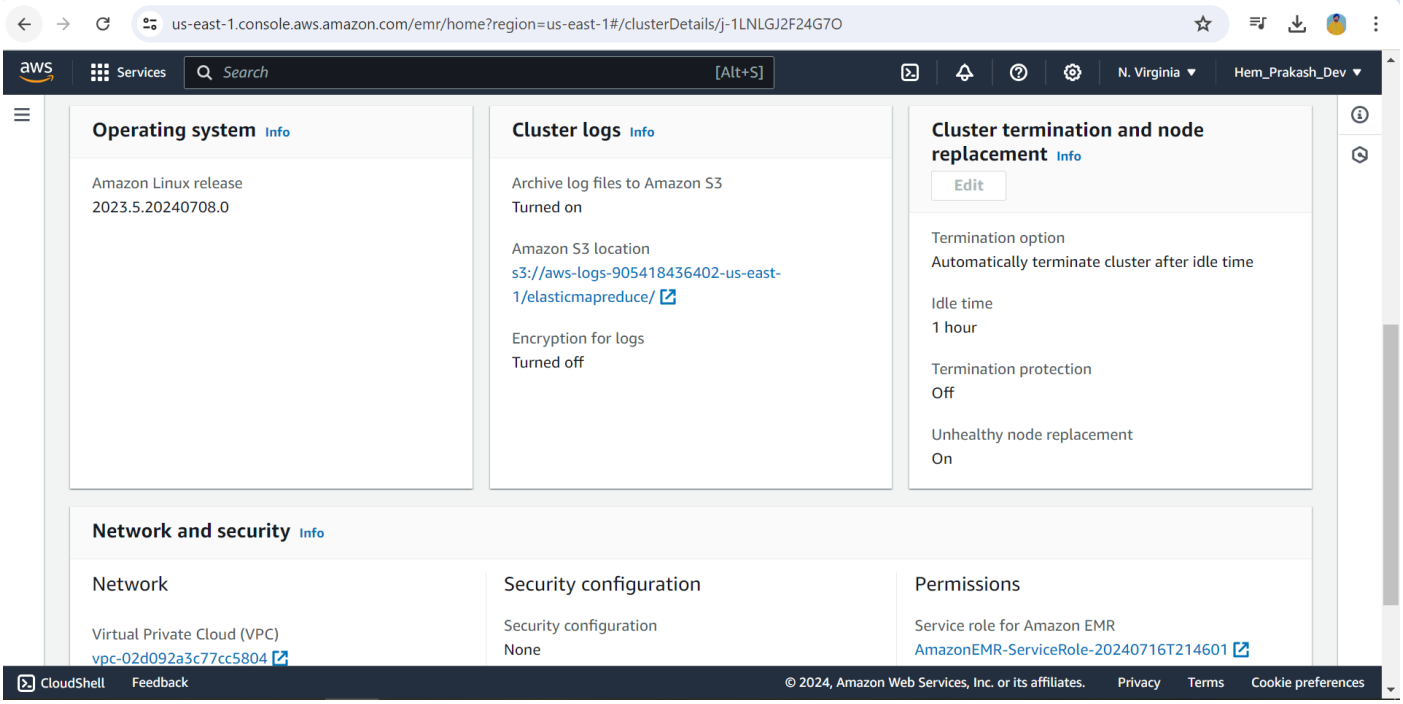
us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-1LNLGJ2F24G7O
Services
Search
[Alt+S]
N. Virginia
Hem_Prakash_Dev
Operating system Info
Amazon Linux release
2023.5.20240708.0
Cluster logs Info
Archive log files to Amazon S3
Turned on
Amazon S3 location
s3://aws-logs-905418436402-us-east-1/elasticmapreduce/
Encryption for logs
Turned off
Cluster termination and node replacement Info
Edit
Termination option
Automatically terminate cluster after idle time
Idle time
1 hour
Termination protection
Off
Unhealthy node replacement
On
Network and security Info
Network
Virtual Private Cloud (VPC)
vpc-02d092a3c77cc5804
Security configuration
Security configuration
None
Permissions
Service role for Amazon EMR
AmazonEMR-ServiceRole-20240716T214601
CloudShell
Feedback
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

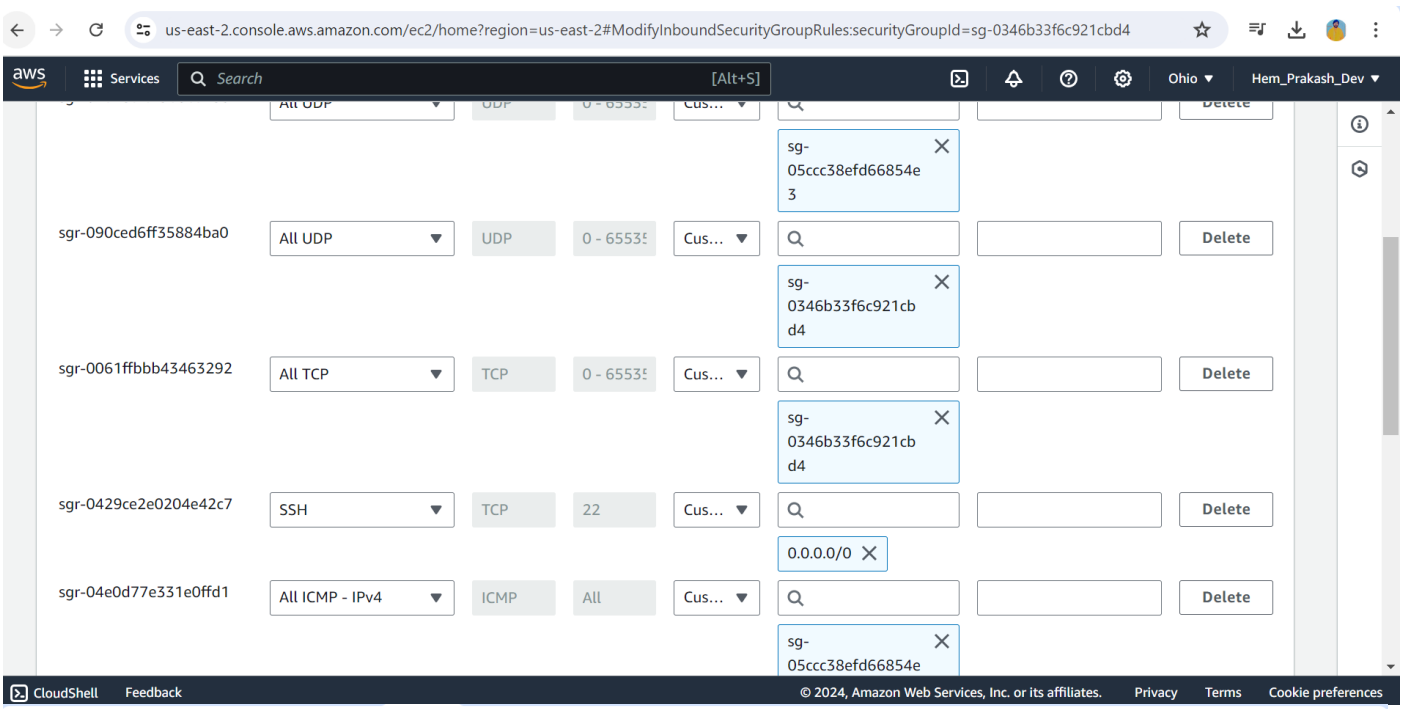
us-east-2.console.aws.amazon.com/ec2/home?region=us-east-2#ModifyInboundSecurityGroupRules:securityGroupId=sg-0346b33f6c921cbd4
aws
Services
Search
[Alt+S]
Ohio
Hem_Prakash_Dev
sg-05ccc38efd66854e3
sgr-090ced6ff35884ba0
All UDP
UDP
0 - 65535
Cus...
sg-0346b33f6c921cbd4
Delete
sgr-0061ffbbb43463292
All TCP
TCP
0 - 65535
Cus...
sg-0346b33f6c921cbd4
Delete
sgr-0429ce2e0204e42c7
SSH
TCP
22
Cus...
0.0.0.0/0
Delete
sgr-04e0d77e331e0ffd1
All ICMP - IPv4
ICMP
All
Cus...
sg-05ccc38efd66854e
Delete
CloudShell
Feedback
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

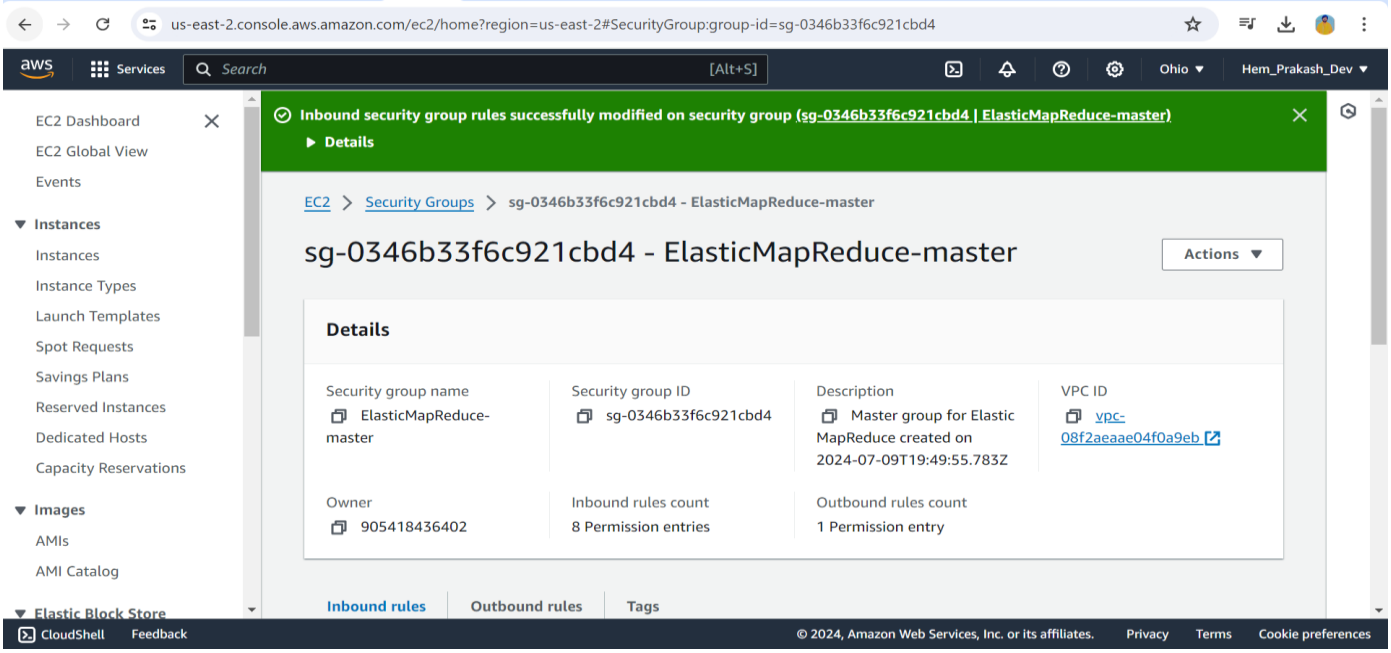
us-east-2.console.aws.amazon.com/ec2/home?region=us-east-2#SecurityGroup:group-id=sg-0346b33f6c921cbd4
aws
Services
Search
[Alt+S]
Ohio
Hem_Prakash_Dev
EC2 Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Inbound security group rules successfully modified on security group (sg-0346b33f6c921cbd4 | ElasticMapReduce-master)
Details
EC2 > Security Groups > sg-0346b33f6c921cbd4 - ElasticMapReduce-master
sg-0346b33f6c921cbd4 - ElasticMapReduce-master
Actions
Details
Security group name
ElasticMapReduce-master
Security group ID
sg-0346b33f6c921cbd4
Description
Master group for Elastic MapReduce created on 2024-07-09T19:49:55.783Z
VPC ID
vpc-08f2aeaae04f0a9eb
Owner
905418436402
Inbound rules count
8 Permission entries
Outbound rules count
1 Permission entry
Inbound rules
Outbound rules
Tags
CloudShell
Feedback
© 2024, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

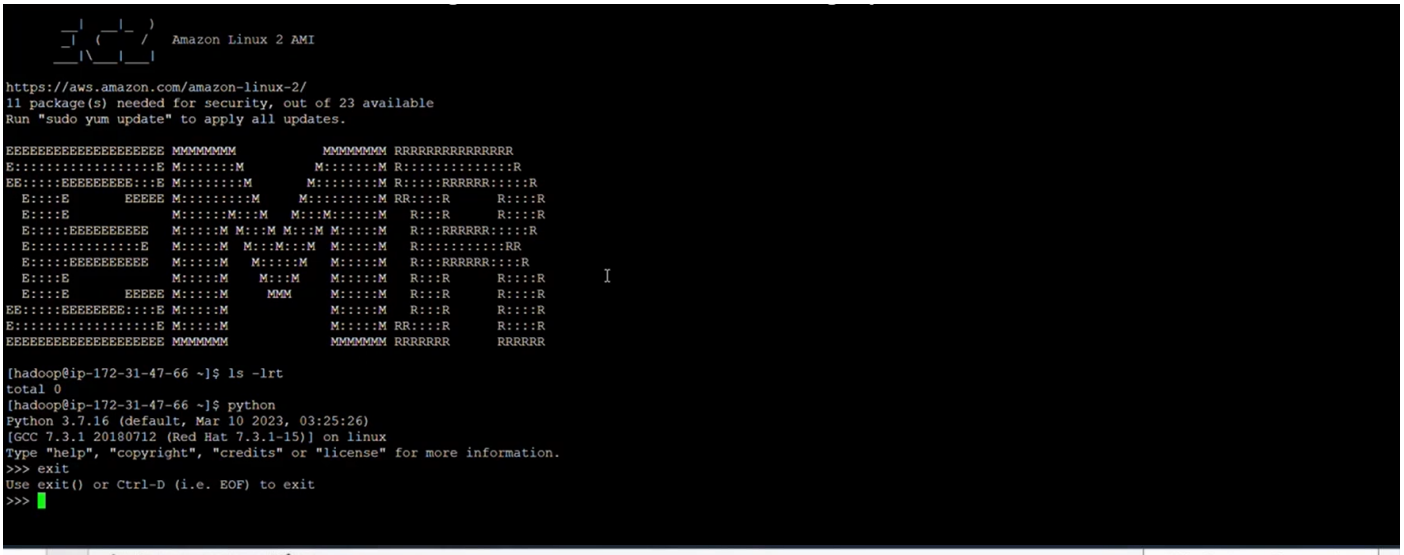

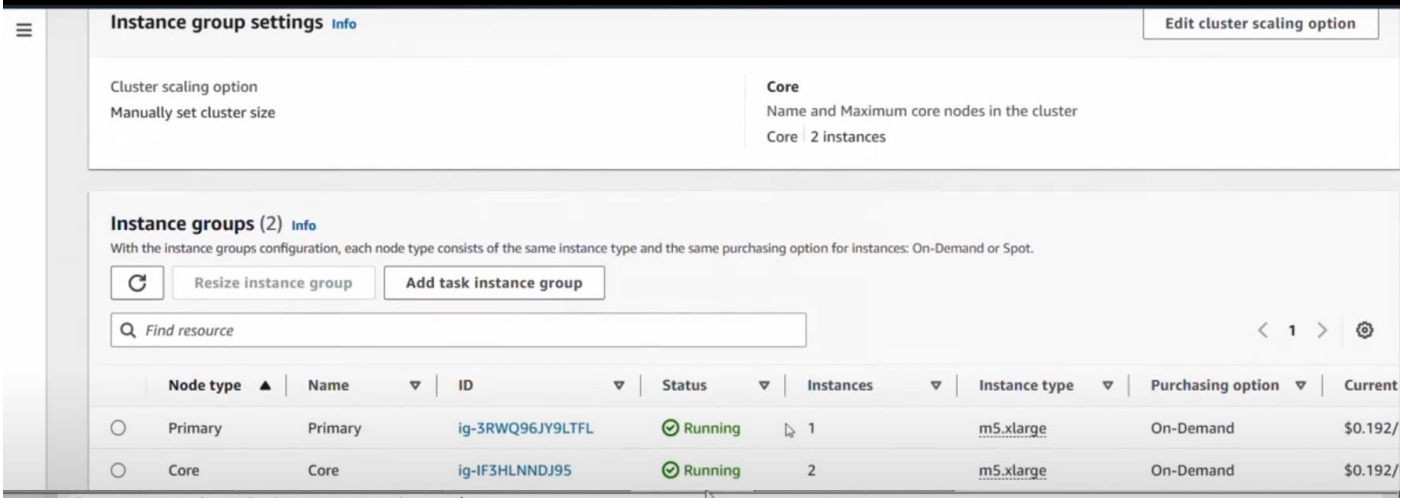

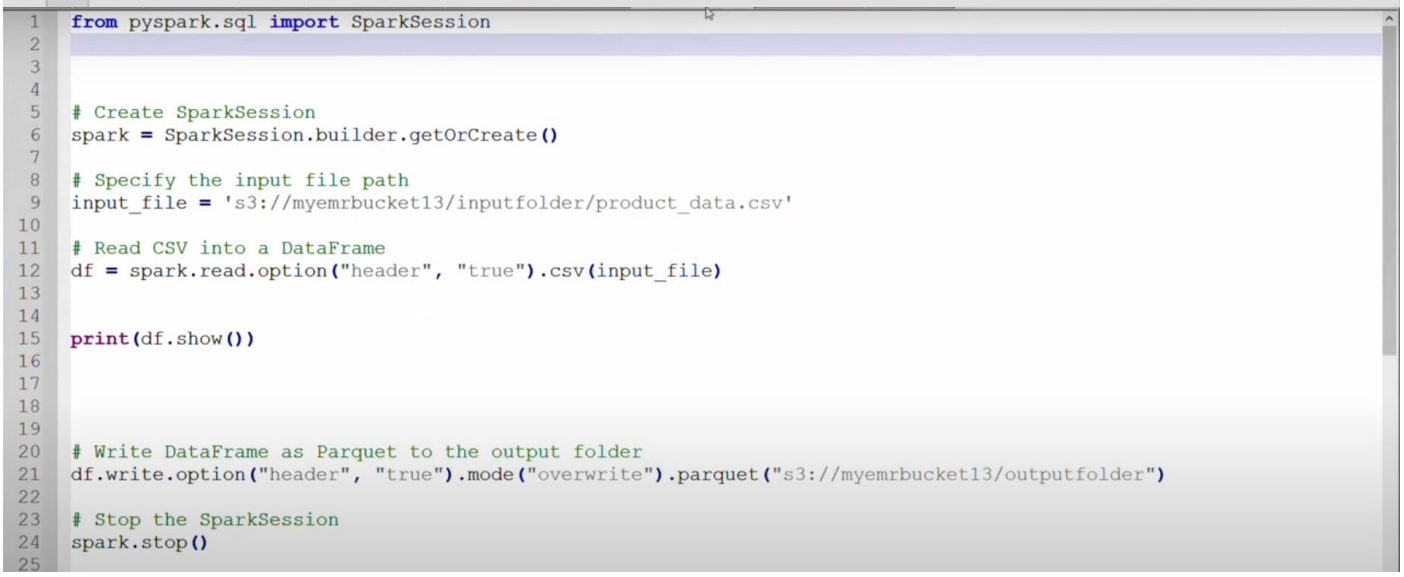

```
from pyspark.sql import SparkSession



# Create SparkSession
spark = SparkSession.builder.getOrCreate()

# Specify the input file path
input_file = 's3://myemrbucket13/inputfolder/product_data.csv'

# Read CSV into a DataFrame
df = spark.read.option("header", "true").csv(input_file)


print(df.show())




# Write DataFrame as Parquet to the output folder
df.write.option("header", "true").mode("overwrite").parquet("s3://myemrbucket13/outputfolder")

# Stop the SparkSession
spark.stop()
```

3. Apply data transformation and cleansing techniques to prepare the data for aggregation and analysis.

Applying data transformation and cleansing techniques to prepare the data for aggregation and analysis.

- **Transformation**: Filter, select, join, and aggregate data as needed.
- **Cleansing**: Handle missing values, remove duplicates, and standardize formats.

```python
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, trim

# Initialize a SparkSession
spark = SparkSession.builder.appName("DataCleaningExample").getOrCreate()

# Sample data
data = [
    (1, "  Alice  ", None),
    (2, "Bob", "value2"),
    (3, "Alice", "value3"),
    (1, "  Alice  ", "value1"),
    (4, None, "value4")
]

# Column names
columns = ["id", "name", "column_name"]

# Create DataFrame
df = spark.createDataFrame(data, columns)

# Remove duplicates
df_cleaned = df.dropDuplicates()

# Handle missing values
df_cleaned = df_cleaned.na.fill({'column_name': 'default_value', 'name': 'unknown'})

# Standardize formats
df_cleaned = df_cleaned.withColumn('trimmed_column', trim(col('name')))

# Show the result
df_cleaned.show()

# Stop the SparkSession
spark.stop()
```

Output of the sample taken

```
+---+---------+-------------+--------------+
| id|     name|  column_name|trimmed_column|
+---+---------+-------------+--------------+
|  1|   Alice |   value1    |         Alice|
|  1|   Alice |      null   |         Alice|
|  3|    Alice|   value3    |         Alice|
|  2|      Bob|    value2   |           Bob|
|  4|   unknown|   value4   |       unknown|
+---+---------+-------------+--------------+
```

4. Implement data partitioning and indexing strategies to optimize query performance.

Implementing data partitioning and indexing strategies to optimize query performance.

Partitioning: Partition data based on commonly queried fields.

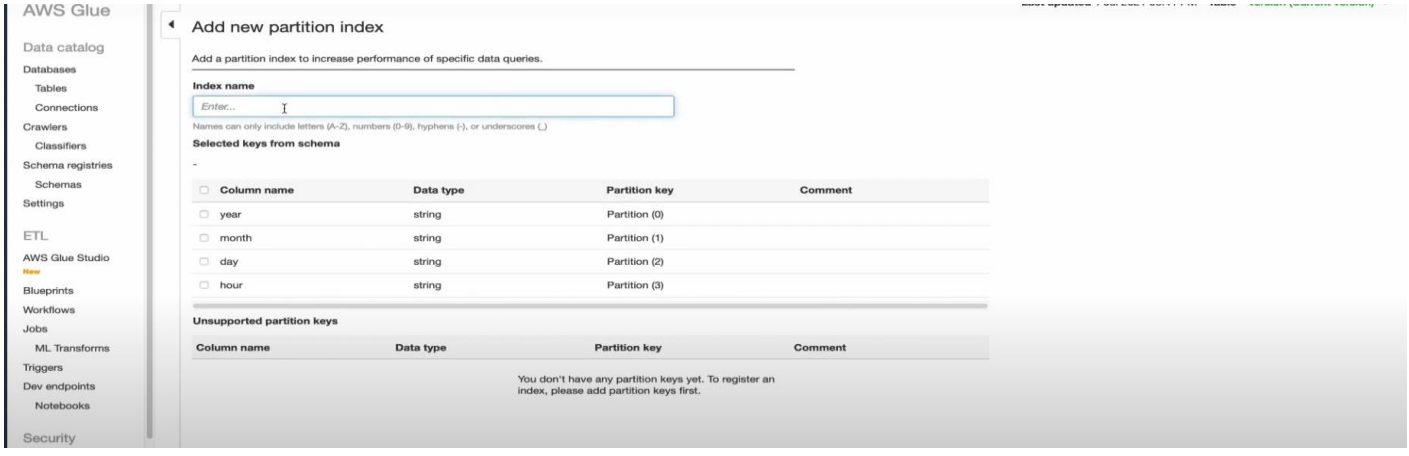

```python
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, trim

# Initialize a SparkSession with Hadoop AWS package
spark = SparkSession.builder \
    .appName("DataCleaningExample") \
    .config("spark.hadoop.fs.s3a.impl", "org.apache.hadoop.fs.s3a.S3AFileSystem") \
    .config("spark.hadoop.fs.s3a.aws.credentials.provider", "com.amazonaws.auth.DefaultAW
    .getOrCreate()

# Sample data
data = [
    (1, "  Alice  ", None, "2024-01-01"),
    (2, "Bob", "value2", "2024-01-02"),
    (3, "Alice", "value3", "2024-01-01"),
    (1, "  Alice  ", "value1", "2024-01-03"),
    (4, None, "value4", "2024-01-02")
]

# Column names
columns = ["id", "name", "column_name", "partition_column"]

# Create DataFrame
df = spark.createDataFrame(data, columns)

# Remove duplicates
df_cleaned = df.dropDuplicates()

# Handle missing values
df_cleaned = df_cleaned.na.fill({'column_name': 'default_value', 'name': 'unknown'})

# Standardize formats
df_cleaned = df_cleaned.withColumn('trimmed_column', trim(col('name')))

# Show the result
df_cleaned.show()

# Write to S3 in JSON format, partitioned by 'partition_column'
df_cleaned.write.partitionBy('partition_column').json('s3://my-first-bucket')

# Stop the SparkSession
spark.stop()
```

```
+---+---------+-------------+--------------+--------------+
| id|     name|  column_name|partition_column|trimmed_column|
+---+---------+-------------+--------------+--------------+
|  1|   Alice |    value1   |    2024-01-03 |      Alice   |
|  1|   Alice |     null    |    2024-01-01 |      Alice   |
|  3|    Alice|    value3   |    2024-01-01 |      Alice   |
|  2|      Bob|    value2   |    2024-01-02 |       Bob    |
|  4|  unknown|    value4   |    2024-01-02 |   unknown    |
+---+---------+-------------+--------------+--------------+
```

5. Update the GitHub repository with the code and configuration files for data ingestion and processing.

Updating the GitHub repository with the code and configuration files for data ingestion and processing.

1. **Initializing a Git Repository**:
   git init
2. **Adding and Commit Code**:
   git add .
   git commit -m "Initial commit with data ingestion and processing scripts"
3. **Pushing to GitHub:**

git remote add origin https://github.com/Hem-Prakash-Dev-Bharadwaj/data-engineering
git push -u origin main